

1. Installing dependencies and setup

```
In [ ]: !pip install --upgrade tensorflow[and-cuda] opencv-python matplotlib
```

2. Checking installations

```
In [ ]: !pip list
```

3. Importing libraries

```
In [ ]: import os
os.environ["CUDA_VISIBLE_DEVICES"] = "-1" # Disables GPU
import tensorflow as tf
import cv2
import imghdr
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy
import cv2
from tensorflow.keras.models import load_model
```

4. Removing dodgy images from data sets

```
In [ ]: data_dir = 'data'
```

```
In [ ]: image_exts = ['.jpeg', '.jpg', '.bmp', '.png']
```

5. Loading data

```
In [ ]: data = tf.keras.utils.image_dataset_from_directory('data')
```

```
In [ ]: data_iterator = data.as_numpy_iterator()
```

```
In [ ]: batch = data_iterator.next()
```

```
In [ ]: #checking what is assigned to 0 and 1

fig, ax = plt.subplots(ncols=4, figsize=(20,20))
for idx, img in enumerate(batch[0][:4]):
    ax[idx].imshow(img.astype(int))
```

```
ax[idx].title.set_text(batch[1][idx])

#happy = 0
#sad = 1
```

6. Scaling Data from 0-255 to 0-1

```
In [ ]: data = data.map(lambda x,y: (x/255, y))
```

```
In [ ]: data.as_numpy_iterator().next()
```

7. Splitting data

```
In [ ]: len(data)
```

```
In [ ]: train_size = int(len(data)*.8)
        val_size = int(len(data)*.1)
        test_size = int(len(data)*.1)
```

```
In [ ]: print(train_size)
        print(val_size)
        print(test_size)
```

```
In [ ]: train = data.take(train_size) #take 8 batches
        val = data.skip(train_size).take(val_size) #skip batches already taken and take 1 b
        test = data.skip(train_size+val_size).take(test_size) #skip batches already taken a
```

8. Building Model

```
In [ ]: model = Sequential()
```

```
In [ ]: tf.random.set_seed(1234)

model = Sequential(
    [
        ### START CODE HERE ###
        Conv2D(16, (3,3), strides=1, activation='relu', input_shape=(256,256,3), na
        MaxPooling2D(name="maxpool1"),

        Conv2D(32, (3,3), strides=1, activation='relu', name="conv2"),
        MaxPooling2D(name="maxpool2"),

        Conv2D(16, (3,3), strides=1, activation='relu', name="conv3"),
        MaxPooling2D(name="maxpool3"),

        Flatten(name="flatten"),
        Dense(256, activation='relu', name="dense1"),
        Dense(1, activation='sigmoid', name="output")
```

```
    ], name="my_model"
)
```

```
In [ ]: model.compile('adam', loss=tf.losses.BinaryCrossentropy(), metrics=['accuracy'])
```

```
In [ ]: model.summary()
```

9. Training

```
In [ ]: logdir='logs'
```

```
In [ ]: tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir)
```

```
In [ ]: hist = model.fit(train, epochs=20, validation_data=val, callbacks=[tensorboard_callback])
```

10. Plotting

```
In [ ]: fig = plt.figure()
plt.plot(hist.history['loss'], color='blue', label='loss')
plt.plot(hist.history['val_loss'], color='red', label='val_loss')
fig.suptitle('Loss', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```

```
In [ ]: fig = plt.figure()
plt.plot(hist.history['accuracy'], color='blue', label='accuracy')
plt.plot(hist.history['val_accuracy'], color='red', label='val_accuracy')
fig.suptitle('Accuracy', fontsize=20)
plt.legend(loc="upper left")
plt.show()
```

11. Evaluating Performance

```
In [303... precision = Precision()
recall = Recall()
accuracy = BinaryAccuracy()

# Ensure dataset iteration does not go out of range
for batch in test.take(len(test)): # Use `.take()` to iterate safely
    X, y = batch
    yhat = model.predict(X)

    precision.update_state(y, yhat)
    recall.update_state(y, yhat)
    accuracy.update_state(y, yhat)

print(f'Precision: {precision.result().numpy()}, Recall: {recall.result().numpy()},
```

```
print("yhat shape:", yhat.shape)
```

2025-01-20 22:37:04.601705: W tensorflow/core/lib/png/png_io.cc:89] PNG warning: iCC P: known incorrect sRGB profile

1/1 ————— 0s 198ms/step
Precision: 1.0, Recall: 1.0, Accuracy: 1.0
yhat shape: (32, 1)

```
In [304... img = cv2.imread('happyExample2.jpg')
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.show()
```



```
In [ ]: resizedImg = tf.image.resize(img, (256,256))
plt.imshow(resizedImg.numpy().astype(int))
plt.show()
```

```
In [ ]: np.expand_dims(resizedImg, 0)
```

```
In [307... yhat = model.predict(np.expand_dims(resizedImg/255, 0))
print(yhat)

if yhat > 0.5:
    print("Prediction: Sad")
else:
    print("Prediction: Happy")
```

1/1 ————— 0s 102ms/step
[[0.01000653]]
Prediction: Happy

```
In [308... img = cv2.imread('sadExample1.jpg')
plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
```

```
plt.show()
```



```
In [ ]: resizedImg = tf.image.resize(img, (256,256))
plt.imshow(resizedImg.numpy().astype(int))
plt.show()
```

```
In [ ]: np.expand_dims(resizedImg, 0)
```

```
In [311... yhat = model.predict(np.expand_dims(resizedImg/255, 0))
print(yhat)

if yhat > 0.5:
    print("Prediction: Sad")
else:
    print("Prediction: Happy")
```

1/1 ————— 0s 67ms/step

[[0.9501678]]

Prediction: Sad

12. Saving the Model

```
In [ ]: model.save(os.path.join('models', 'EmotionDetector.h5'))
```

```
In [314... new_model = load_model(os.path.join('models', 'EmotionDetector.h5'))
yhatnew = new_model.predict(np.expand_dims(resizedImg/255, 0))
print(yhatnew)

if yhatnew > 0.5:
    print("Prediction: Sad")
```

```
else:  
    print("Prediction: Happy")
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

WARNING:tensorflow:6 out of the last 13 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7f21e9a02cb0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

WARNING:tensorflow:6 out of the last 13 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7f21e9a02cb0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/1  0s 155ms/step

[[0.9501678]]

Prediction: Sad